

# Ensemble Learning

Business Analytics  
Stefan Feuerriegel

# Outline

- 1 Decision Trees
- 2 Random Forests
- 3 Boosting
- 4 AdaBoost

# Setup

- ▶ Accessing credit scores

```
library(caret)  
data(GermanCredit)
```

- ▶ Split data into index subset for **training** (20 %) and **testing** (80 %) instances

```
inTrain <- runif(nrow(GermanCredit)) < 0.2
```

# Outline

- 1** Decision Trees
- 2 Random Forests
- 3 Boosting
- 4 AdaBoost

# Decision Trees in R

- ▶ Loading required libraries `rpart`, `party` and `partykit`

```
library(rpart)
library(party)
library(partykit)
```

- ▶ Building a decision tree with

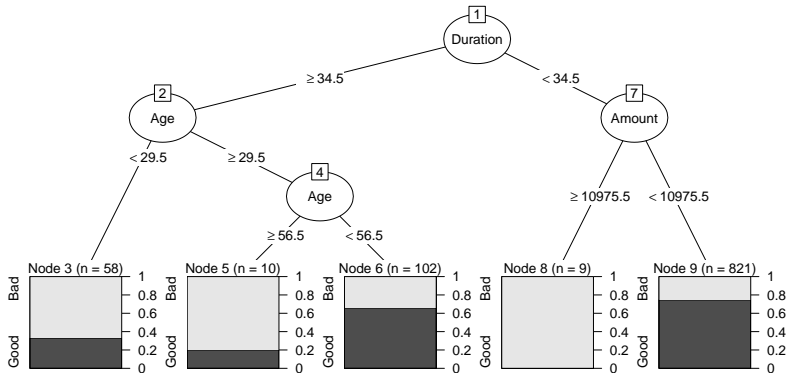
```
rpart(formula, method="class", data=d)
```

```
dt <- rpart(Class ~ Duration + Amount + Age,
             method="class", data=GermanCredit)
```

# Decision Trees in R

- Plot decision tree using `plot(dt)`

```
plot(as.party(dt))
```



# Prediction with Decision Trees

- ▶ `predict(dt, test, type="class")` predicts classes on new data test

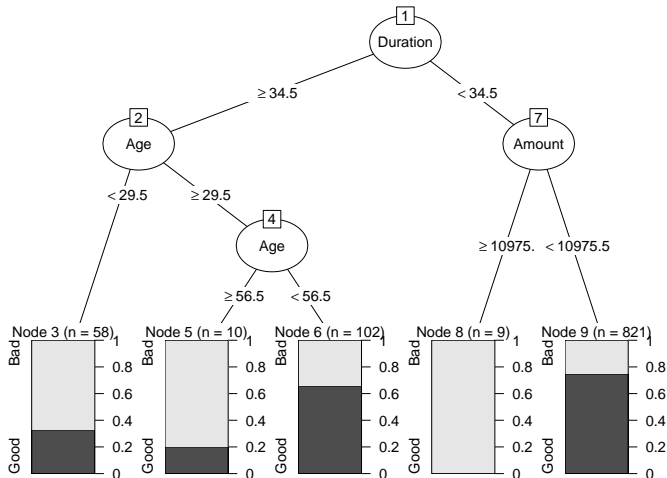
```
pred <- predict(dt, GermanCredit, type="class")
pred[1:5]

##      1      2      3      4      5
## Good  Bad Good Good Good
## Levels: Bad Good
```

- ▶ Output: **predicted label in 1st row** out of all possible labels (2nd row)
- ▶ Pruning occurs through `prune(dt, cp = ...)` with a given complexity parameter
  - ▶ Usual heuristic:  
`dt$cptable[which.min(dt$cptable[, "xerror"]), "CP"]`

# Pruning Decision Trees

```
p <- prune(dt, cp = dt$cptable[which.min(dt$cptable[, "xerror"]), "CP"])\nplot(as.party(p))
```





# Outline

1 Decision Trees

**2 Random Forests**

3 Boosting

4 AdaBoost

# Random Forests in R

- ▶ Load required library `randomForest`

```
library(randomForest)
```

- ▶ Learn random forest on training data with `randomForest(...)`

```
rf <- randomForest(Class ~ .,  
                    data=GermanCredit,  
                    ntree=100)
```

- ▶ Options to control behavior
  - ▶ `ntree` controls the number of trees (default: 500)
  - ▶ `mtry` gives number of variables to choose from at each node
  - ▶ `na.action` specifies how to handle missing values
  - ▶ `importance=TRUE` calculates variable importance metric
- ▶ Predict credit scores for testing instances

```
pred <- predict(rf, newdata=GermanCredit)
```

# Variable Importance in R


- Learn random forest and enable the calculation of variable importance metrics via `importance=TRUE`

```
rf2 <- randomForest(Class ~ .,  
                     data=GermanCredit, #with full dataset  
                     ntree=100,  
                     importance=TRUE)
```

# Variable Importance in R

- Plot variable importance via `varImpPlot(rf, ...)`

```
varImpPlot(rf2, type=1, n.var=5)
```



- `type` choose the importance metric (= 1 is the mean decrease in accuracy if the variable would be randomly permuted)
- `n.var` denotes number of variables

# Outline

- 1 Decision Trees
- 2 Random Forests
- 3 Boosting**
- 4 AdaBoost

# Boosting in R

- Load the required packages `mboost`

```
library(mboost)
```

- Fit a **generalized linear model** via `glmboost(...)`

```
m.boost <- glmboost(Class ~ Amount + Duration
                    + Personal.Female.Single,
                    family=Binomial(), # needed for classification
                    data=GermanCredit)

coef(m.boost)

##      (Intercept)      Amount      Duration
## 4.104949e-01 -1.144369e-05 -1.703911e-02
## attr(,"offset")
## [1] 0.4236489
```

- Different from the normal `glm(...)` routine, the boosted version inherently **performs variable selection**

# Outline

- 1 Decision Trees
- 2 Random Forests
- 3 Boosting
- 4 AdaBoost**

# AdaBoost in R

- Load required package `ada`

```
library(ada)
```

- Fit AdaBoost model on training data with `ada(..., iter)` given a fixed number `iter` of iterations

```
m.ada <- ada(Class ~ .,  
              data=GermanCredit,  
              iter=50)
```

- Evaluate on test data

```
pred <- predict(m.ada, newdata=GermanCredit)
```